

Exercises on MapReduce

Exercise 1 Design an MR algorithm which solves (exactly) the Maximum pairwise distance problem for a set S of N points from a metric space (M, d) . The algorithm must run in $R = O(1)$ rounds and require local space $M_L = O(\sqrt{N})$ and aggregate space $M_A = O(N^2)$.

Solution. The input-output specification of the problem is the following one:

- **Input:** Set S of N points represented by pairs (i, x_i) , for $0 \leq i < N$.
- **Output:** $(0, d_{\max} = \max_{0 \leq i, j < N} d(x_i, x_j))$.

The idea is first to compute the distances between all possible pairs x_i, x_j with $0 \leq i \leq j < N$ (this is done in the first 2 rounds), and then to compute the maximum in several rounds, processing distances in groups of size \sqrt{N} . The details of the algorithm are below.

Round 1

- *Map phase:* map each pair (i, x_i) into the following \sqrt{N} intermediate pairs: $((i, j), x_i)$, with $0 \leq j < \sqrt{N}$.
- *Reduce phase:* empty

Round 2

- *Map phase:* map each pair $((i, j), x_i)$ into the following \sqrt{N} intermediate pairs: $((i, j * \sqrt{N} + j'), x_i)$, with $0 \leq j' < \sqrt{N}$.
- *Reduce phase:* empty

Observation: At this point, for every x_i we have N pairs $((i, j), x_i)$, for $0 \leq i < N$.

Round 3

- *Map phase:* map each pair $((i, j), x_i)$ into the pair $((\min\{i, j\}, \max\{i, j\}), x_i)$.
- *Reduce phase:* for every $0 \leq i \leq j < N$ gather all pairs with key (i, j) . Note that if $i < j$ there will be exactly two such pairs: $((i, j), x_i)$ and $((i, j), x_j)$ (otherwise there will be only one, $((i, i), x_i)$). If $i < j$, produce the output pair $(i * N + j, d(x_i, x_j))$.

Observation: At this point we have all pairwise distances among points. Note also, that Rounds 2 and 3 can be easily merged into one round, but we kept them separate for clarity.

Round 4

- *Map phase*: map each pair (k, d) into the intermediate pair $(k' = k \bmod N^{3/2}, d)$ and note that $k' \in [0, N^{3/2})$.
- *Reduce phase*: for each $k' \in [0, N^{3/2})$ independently, gather the $\leq \sqrt{N}$ intermediate pairs (k', d) and produce the pair (k', d') where d' is the maximum of the values of all such intermediate pairs.

Round 5

- *Map phase*: map each pair (k, d) into the intermediate pair $(k' = k \bmod N, d)$ and note that $k' \in [0, N)$.
- *Reduce phase*: for each $k' \in [0, N)$ independently, gather the $\leq \sqrt{N}$ intermediate pairs (k', d) and produce the pair (k', d') where d' is the maximum of the values of all such intermediate pairs.

Round 6

- *Map phase*: map each pair (k, d) into the pair $(k' = k \bmod \sqrt{N}, d)$ and note that $k' \in [0, \sqrt{N})$.
- *Reduce phase*: for each $k' \in [0, \sqrt{N})$ independently, gather the $\leq \sqrt{N}$ intermediate pairs (k', d) and produce the pair $(0, d')$ where d' is the maximum of the values of all such intermediate pairs.

Round 7

- *Map phase*: identity map.
- *Reduce phase*: gather the $\leq \sqrt{N}$ pairs $(0, d)$ and return, as output, the pair with maximum value.

It is easy to see that the above 7-round algorithm uses aggregate space proportional to the number of ordered pairs of points, that is, $M_A = O(N^2)$. As for the local space, in each round the map phase transforms a single pair into at most \sqrt{N} intermediate pairs, while the reduce phase works separately on subsets of size $\Theta(\sqrt{N})$ and produces one pair out of each subset. Hence, the local space M_L is $\Theta(\sqrt{N})$. \square

Exercise 1b *Design an MR algorithm which solves (exactly) the Maximum pairwise distance problem for a set S of N points from a metric space (M, d) . The algorithm must run in $R = \Theta(\sqrt{N})$ rounds and require local space $M_L = \Theta(\sqrt{N})$ and aggregate space $M_A = \Theta(N)$.*

Solution. The input-output specification is the same as in the previous exercise. The idea now is to partition the N pairs into \sqrt{N} sets $S_0, S_1, \dots, S_{\sqrt{N}-1}$ of size \sqrt{N} , and then to search for the maximum pairwise distance by searching within the sets $S_i \cup S_j$ for each

$0 \leq i, j < \sqrt{N}$. By using \sqrt{N} rounds, we compute the N subproblems by solving \sqrt{N} of them per round.

More in details, we first compute the maximum pairwise distance for each set S_i : in round j for $0 \leq j < \sqrt{N}$, we compute the maximum pairwise distance from points in S_i to points in $S_{(i+j) \bmod \sqrt{N}}$; if the distance is larger than the maximum distance found in the previous rounds, we keep it; otherwise, we disregards it. During the last round \sqrt{N} , we receive the distances computed for each S_i and we emit the maximum pairwise distances.

The details of the algorithm are below (we assume that the input of each round $0 \leq i < \sqrt{N}$ consists of the N input pairs (i, x_i) and the \sqrt{N} pairs emitted by the previous round, if any).

Round 1

- **Observation:** In this round, we compute the maximum pairwise distance within each set S_i .
- *Map phase:* map each pair (i, x_i) into the pair $(i \bmod \sqrt{N}, x_i)$.
- *Reduce phase:* For every key $0 \leq k < \sqrt{N}$, independently, gather the \sqrt{N} pairs (k, x) in the set S_k , compute $d_{\max, k} = \max_{x, y \in S'_k} d(x, y)$ and then emit $(-1, (k, d_{\max, k}))$. (The key value -1 is to distinguish, in the following map, the pairs with $d_{\max, k}$ values from the input pairs.

Round i , for each $1 \leq i < \sqrt{N}$

- **Observation:** In this round, we compute the maximum pairwise distance from points in S_i to points in $S_{i+j \bmod \sqrt{N}}$.
- *Map phase:* map each pair (i, x_i) , with $i \geq 0$, into the two pairs $(i \bmod \sqrt{N}, x_i)$ and $((j + i) \bmod \sqrt{N}, x_i)$. Map each pair $(-1, (k, d_k))$ into $(k, (k, d_k))$ (these are the \sqrt{N} pairs containing the partial maximum pairwise distance for each set S_i).
- *Reduce phase:* For every key $0 \leq k < \sqrt{N}$, independently, collect the pair $(k, (k, d_k))$ and gather the $2\sqrt{N}$ pairs (k, x) in the set S'_k ; then compute $d_{\max, k} = \max_{x, y \in S'_k} d(x, y)$ and then emit $(-1, (k, \max\{d_k, d_{\max, k}\}))$.

Round \sqrt{N}

- **Observation:** This round computes the max among all the partial maximum distances.
- *Map phase:* Map each pair $(-1, (k, d_k))$ into $(0, (k, d_k))$.
- *Reduce phase:* Gather the \sqrt{N} intermediate pairs $(0, d_k)$, compute $d_{\max} = \max_{0 \leq k < \sqrt{N}} d_k$ and produce the pair $(0, d_{\max})$.

Since there are two copies of each input pair in each round and \sqrt{N} pairs with the partial maximum pairwise distances, the above $(\sqrt{N} + 1)$ -round algorithm uses linear aggregate space, that is, $M_A = \Theta(N)$. As for the local space, in each round the map phase splits the input pairs into sets of \sqrt{N} pairs; hence, the local space M_L is $\Theta(N^{1/2})$.

□

Exercise 2 (Exercise 2.3.1.(b) of [LRU14]) *Design an $O(1)$ -round MapReduce algorithm to compute the arithmetic mean of a set S of N integers, using $O(\sqrt{N})$ local space and linear aggregate space. Assume that S is provided in input as a set of key-value pairs (i, x_i) where x_i is an arbitrary integer and i , with $0 \leq i < N$, is the key of the pair.*

Solution. The input-output specification of the problem is the following one:

- **Input:** $S = \{(i, x_i) : 0 \leq i < N\}$.
- **Output:** $(0, \text{mean}(S))$, where $\text{mean}(S) = (1/N) \sum_{i=0}^{N-1} x_i$.

While the trivial 1-round algorithm uses $\Theta(N)$ local space, the space requirements can be substantially lowered using the following simple 2-round strategy. (We assume that the input size N is known to the algorithm.)

Round 1

- *Map phase:* map each pair (i, x_i) into the intermediate pair $(i \bmod \sqrt{N}, x_i)$ (for simplicity, assume that \sqrt{N} is an integer).
- *Reduce phase:* for each key k independently, with $0 \leq k < \sqrt{N}$, gather the set S_k of all intermediate pairs (k, x) and produce the pair $(0, \text{sum}_k)$, where $\text{sum}_k = \sum_{(k, x) \in S_k} x$.

Round 2

- *Map phase:* empty.
- *Reduce phase:* gather all pairs $(0, \text{sum}_k)$, with $0 \leq k < \sqrt{N}$, and return the output pair $(0, (1/N) \sum_{k=0}^{\sqrt{N}-1} \text{sum}_k)$.

It is immediate to see that the algorithm requires $M_L = \Theta(\sqrt{N})$ and aggregate space $M_A = \Theta(N)$. □

Exercise 3 (Exercise 2.3.1.(d) of [LRU14]) *Design an $O(1)$ -round MapReduce algorithm to compute the number of distinct integers in a set S of N integers, using $O(\sqrt{N})$ local space and linear aggregate space. Assume that S is provided in input as a set of key-value pairs (i, x_i) where x_i is an arbitrary integer and i , with $0 \leq i < N$, is the key of the pair.*

Solution. The input-output specification of the problem is the following one:

- **Input:** $S = \{(i, x_i) : 0 \leq i < N\}$.
- **Output:** $(0, \text{count})$, where count is the number of distinct integers in S .

We present a straightforward 2-round algorithm and a more space-efficient 4-round algorithm. The 2-round algorithm does a simple removal of duplicates as follows:

Round 1

- *Map phase:* map each pair (i, x_i) into the intermediate pair (x_i, i) , that is, x_i becomes the key.
- *Reduce phase:* for each key x independently, gather the set S_x of all intermediate pairs (x, j) and produce the unique pair $(0, x)$.

Round 2

- *Map phase:* empty.
- *Reduce phase:* gather the set S_0 of all pairs $(0, x)$ and return the output pair $(0, |S_0|)$.

Let N_1 be the maximum number of occurrences of the same integer and N_2 the number of distinct integers in the input set S . Observe that $N_1 \cdot N_2 \geq N$, hence $\max\{N_1, N_2\} \geq \sqrt{N}$. It is immediate to see that the algorithm requires $M_L = \Theta(\max\{N_1, N_2\}) = \Omega(\sqrt{N})$ and aggregate space $M_A = \Theta(N)$. However, note that either N_1 or N_2 , or both, can grow as large as N , asymptotically, hence M_L can be proportional to N in the worst case.

The following alternative strategy reduces the local space to the desired bound. The idea is first to remove all duplicates working on subsets of size at most \sqrt{N} each (Rounds 1 and 2), and then to count the number of surviving, and at this point distinct, integers, counting at most \sqrt{N} of them at once (Rounds 3 and 4)

Round 1

- *Map phase:* map each pair (i, x_i) into the intermediate pair $(i \bmod \sqrt{N}, x_i)$.
- *Reduce phase:* for each key $j \in [0, \sqrt{N})$ independently, gather the set S_j of all intermediate pairs (j, x) and leave one such pair for every distinct integer occurring in S_j .

Observation: after Round 1, for every integer x there are at most \sqrt{N} pairs (j, x) , one for each value of $j \in [0, \sqrt{N})$.

Round 2

- *Map phase:* map each pair (j, x) into the intermediate pair (x, j) , that is, swap key and value.

- *Reduce phase*: for each x independently, gather the set S_x of all intermediate pairs (x, j) and select, arbitrarily, only one such pair.

Observation: after Round 2, for every integer x only one pair survives, and for each value of $j \in [0, \sqrt{N})$ there are at most \sqrt{N} pairs with value j .

Round 3

- *Map phase*: map each pair (x, j) into the intermediate (j, x) , that is swap again key and value.
- *Reduce phase*: for each $j \in [0, \sqrt{N})$ independently, gather all intermediate pairs (j, x) and produce the pair $(0, c_j)$, where c_j is the number of such pairs. (In fact, c_j is a partial count of the distinct integers.

Round 4

- *Map phase*: empty.
- *Reduce phase*: gather all pairs $(0, c_j)$ with $j \in [0, \sqrt{N})$ and return the output pair $(0, \sum_{j=0}^{\sqrt{N}} c_j)$.

It is easy to see that the above 4-round algorithm requires local space $M_L = \Theta(\sqrt{N})$ and aggregate space $M_A = \Theta(N)$. As an example, apply the algorithm to the following input set

$$S = \{(0, 21), (1, 120), (2, 21), (3, 76), (4, 76), (5, 76), (6, 101), (7, 120), \\ (8, 21), (9, 120), (10, 120), (11, 76), (12, 76), (13, 76), (14, 560), (15, 21)\}.$$

□

Exercise 4 Suppose that an instance of the Class count problem is given consisting only of object-class pairs (o, γ) (i.e., without the initial integer keys in $[0, N)$). Consider a modification of the algorithm described in the slides, where in the Map phase of Round 1 each pair (o, γ) is mapped into a pair $(j, (o, \gamma))$, where j is a random integer in $[0, \sqrt{N})$. The rest of the algorithm is unchanged. Analyze probabilistically the local space requirements of the modified algorithm.

Solution. The algorithm to be analyzed is the following simple adaptation of the algorithm presented in Slide 36 on MapReduce.

Round 1

- *Map phase*: map each pair (o_i, γ_i) into the intermediate pair $(x, (o_i, \gamma_i))$, where x (the key of the pair) is a random value in $[0, \sqrt{N})$.

- *Reduce phase*: For each key $x \in [0, \sqrt{N})$ independently, gather the set (say $S^{(x)}$) of all intermediate pairs with key x and, for each distinct class γ labeling some object in $S^{(x)}$, produce the pair $(\gamma, c_x(\gamma))$, where $c_x(\gamma)$ is the number of objects of $S^{(x)}$ labeled with γ .

Round 2

- *Map phase*: empty
- *Reduce phase*: for each class γ independently, gather the at most \sqrt{N} pairs $(\gamma, c_x(\gamma))$ resulting at the end of the previous round, and return the output pair $(\gamma, \sum_x c_x(\gamma))$.

The analysis is similar to the one presented in class for Improved Word count. Let m_x be the number of intermediate pairs with key x produced by the Map phase of Round 1, and let $m = \max_x m_x$. It is easy to see that the above 2-round algorithm requires local space $M_L = O(m + \sqrt{N})$ and aggregate space $M_A = O(N)$. We can prove that $m = O(\sqrt{N})$ with probability $\geq 1 - 1/N^5$ using exactly the same proof that was used to for the theorem on Slides 32-33 on MapReduce. \square

Exercise 5 *Design and analyze an efficient MR algorithm for approximating the maximum pairwise distance, which uses local space $M_L = O(N^{1/4})$. Can you generalize the algorithm to attain $M_L = O(N^\epsilon)$, for any $\epsilon \in (0, 1/2)$?*

Solution. Recall that the input is a set S of N points represented by pairs (i, x_i) , for $0 \leq i < N$, where x_i is the i -th point. The algorithm computes the maximum distance between x_0 and all other points which is within a factor at most 2 from the actual maximum distance between all points, as was seen in class. The input-output specification of the algorithm is the following:

- **Input:** $S = \{(i, x_i) : 0 \leq i < N\}$.
- **Output:** $(0, d_{\max}(0) = \max_{0 \leq i < N} d(x_0, x_i))$.

The algorithm presented next is an adaptation of the one seen in class which complies with the more stringent local space requirements.

Round 1.

- *Map phase*: Pair $(0, x_0)$ is mapped into the set of pairs $\{(j, (0, x_0)) : 0 \leq j < N^{1/4}\}$, while pair (i, x_i) , with $i > 0$, is mapped into pair $(i \bmod N^{3/4}, (i, x_i))$.
- *Reduce phase*: empty.

Round 2.

- *Map phase*: pair $(j, (0, x_0))$ is mapped into the set of pairs $\{(j * N^{1/4} + h, (0, x_0)) : 0 \leq h < N^{1/4}\}$, while all other pairs stay unchanged.

- *Reduce phase*: empty.

Round 3.

- *Map phase*: pair $(j, (0, x_0))$ is mapped into the set of pairs $\{(j * N^{1/4} + h, (0, x_0)) : 0 \leq h < N^{1/4}\}$, while all other pairs stay unchanged.

Observation: after the map phase, there are $N^{3/4}$ copies of x_0 in pairs with indices $j = 0, 1, \dots, N^{3/4} - 1$.

- *Reduce phase*: for each key $0 \leq j < N^{3/4}$ independently, gather $S_j = \{\text{intermediate pairs } (j, (i, x_i))\}$ (it includes $(j, (0, x_0))$), and produce pair $(j, d_{\max}(0, j))$ where $d_{\max}(0, j) = \max$ distance between x_0 and a point in S_j .

Round 4.

- *Map phase*: map each pair (j, d) into $(j \bmod N^{1/2}, d)$.
- *Reduce phase*: for each key $0 \leq j' < N^{1/2}$ independently, gather $S_{j'} = \{\text{intermediate pairs } (j', d)\}$, and produce pair $(j', \max_{(j', d) \in S_{j'}} d)$.

Round 5.

- *Map phase*: map each pair (j, d) into $(j \bmod N^{1/4}, d)$.
- *Reduce phase*: for each key $0 \leq j' < N^{1/4}$ independently, gather $S_{j'} = \{\text{intermediate pairs } (j', d)\}$, and produce pair $(0, \max_{(j', d) \in S_{j'}} d)$.

Round 6.

- *Map phase*: empty
- *Reduce phase*: gather $S_0 = \{\text{all pairs } (0, d)\}$ (at most $N^{1/4}$ pairs) and return the output pair $(0, \max_{(0, d) \in S_0} d)$.

It is easy to see that the above 6-round algorithm uses linear aggregate space. As for the local space, in each round the map phase, if not empty, transforms a single pair into at most $N^{1/4}$ pairs, while the reduce phase, if not empty, works separately on subsets of size $\Theta(N^{1/4})$ and produces one pair out of each subset. Hence, the local space M_L is $\Theta(N^{1/4})$.

The above algorithm can be easily generalized so to require local space $M_L = O(N^\epsilon)$, for any $\epsilon \in (0, 1/2)$. First point x_0 is replicated into $N^{1-\epsilon}$ copies using several rounds, where in each round every copy is replicated into N^ϵ further copies. Then, the input points are partitioned arbitrarily into $N^{1-\epsilon}$ subsets of size N^ϵ each, and a copy of x_0 is gathered together with each subset computing the maximum distance between x_0 and the other points in the subset. Then, in a number of rounds, the maximum of all resulting $N^{1-\epsilon}$ partial max distances is computed, in batches of N^ϵ at a time. It can be seen that $O(1/\epsilon)$ rounds overall are sufficient. \square

Exercise 6 Design an $O(1)$ -round MR algorithm for computing a matrix-vector product $W = A \cdot V$, where A is an $m \times n$ matrix and V is an n -vector. Your algorithm must use $o(n)$ local space and linear, that is, $O(mn)$ aggregate space. For simplicity you may assume $m \leq \sqrt{n}$.

Solution. The input-output specification of the algorithm is as follows. We assume that A is represented through the mn pairs $((i, j), A[i, j])$, with $0 \leq i < m$ and $0 \leq j < n$, where (i, j) is the key, and that V is represented through the n (resp., m) pairs $((j, -1), V[j])$, with $0 \leq j < n$ where $(j, -1)$ is the key. As for W , we assume that in output it will be represented by the pairs $(i, W[i])$, with $0 \leq i < m$.

Note that $W[i]$ is the inner product of a row of A and V , which are both vectors of size n . A trivial approach (1 round) would be to make m copies of V in the map phase, and then, in the reduce phase, to gather each row i of A with a copy of V and compute entry $W[i]$. However, this approach requires $\Theta(n)$ local space. In order to get $o(n)$ local space, we can break V and each row of A into \sqrt{n} segments of size \sqrt{n} each. Then, in the first round, for every row i of A we compute independently the product of each segment of the row and the corresponding segment of V , which yields \sqrt{n} contributions to $W[i]$. In the second round we just add up these \sqrt{n} contributions, separately for every $W[i]$. The following algorithm implements the above idea.

Round 1

- *Map phase*: map each pair $((i, j), A[i, j])$ into pair $((i, s = \lfloor j/\sqrt{n} \rfloor), (i, j, A[i, j]))$, and each pair $((j, -1), V[j])$ into the set of pairs $\{(i, s = \lfloor j/\sqrt{n} \rfloor), (j, -1, V[j])\} : 0 \leq i < m\}$.

Notation: we let $A^{(i,s)}$ denote segment s of row $A^{(i)}$, and let $V^{(i,s)}$ denote segment s of the i -th copy of V .

- *Reduce phase*: for each key (i, s) , $0 \leq i < m$ and $0 \leq s < \sqrt{n}$, independently, gather all pairs with key (i, s) (i.e., $A^{(i,s)}$ and $V^{(i,s)}$), and produce the pair $(i, W_s[i] = A^{(i,s)} \circ V^{(i,s)})$.

Round 2

- *Map phase*: empty
- *Reduce phase*: for each key i independently, gather all pairs $(i, W_s[i])$, $0 \leq s < \sqrt{n}$ and produce the output pair $(i, W[i] = \sum_{s=0}^{\sqrt{n}-1} W_s[i])$.

For what concerns the local space, the map phase of Round 1 works on a single key value pair and transforms such a pair into 1 or $m \leq \sqrt{n}$ intermediate pairs, depending on whether it comes from A or V . The reduce phases of both rounds process sets of $O(\sqrt{n})$ entries each. Thus, we conclude that $M_L \in O(\sqrt{n})$. As for the aggregate space, A is not replicated, while the m replicas of V occupy overall $O(mn)$ space, that is, the same space as A . Thus, $M_A \in O(mn)$. \square